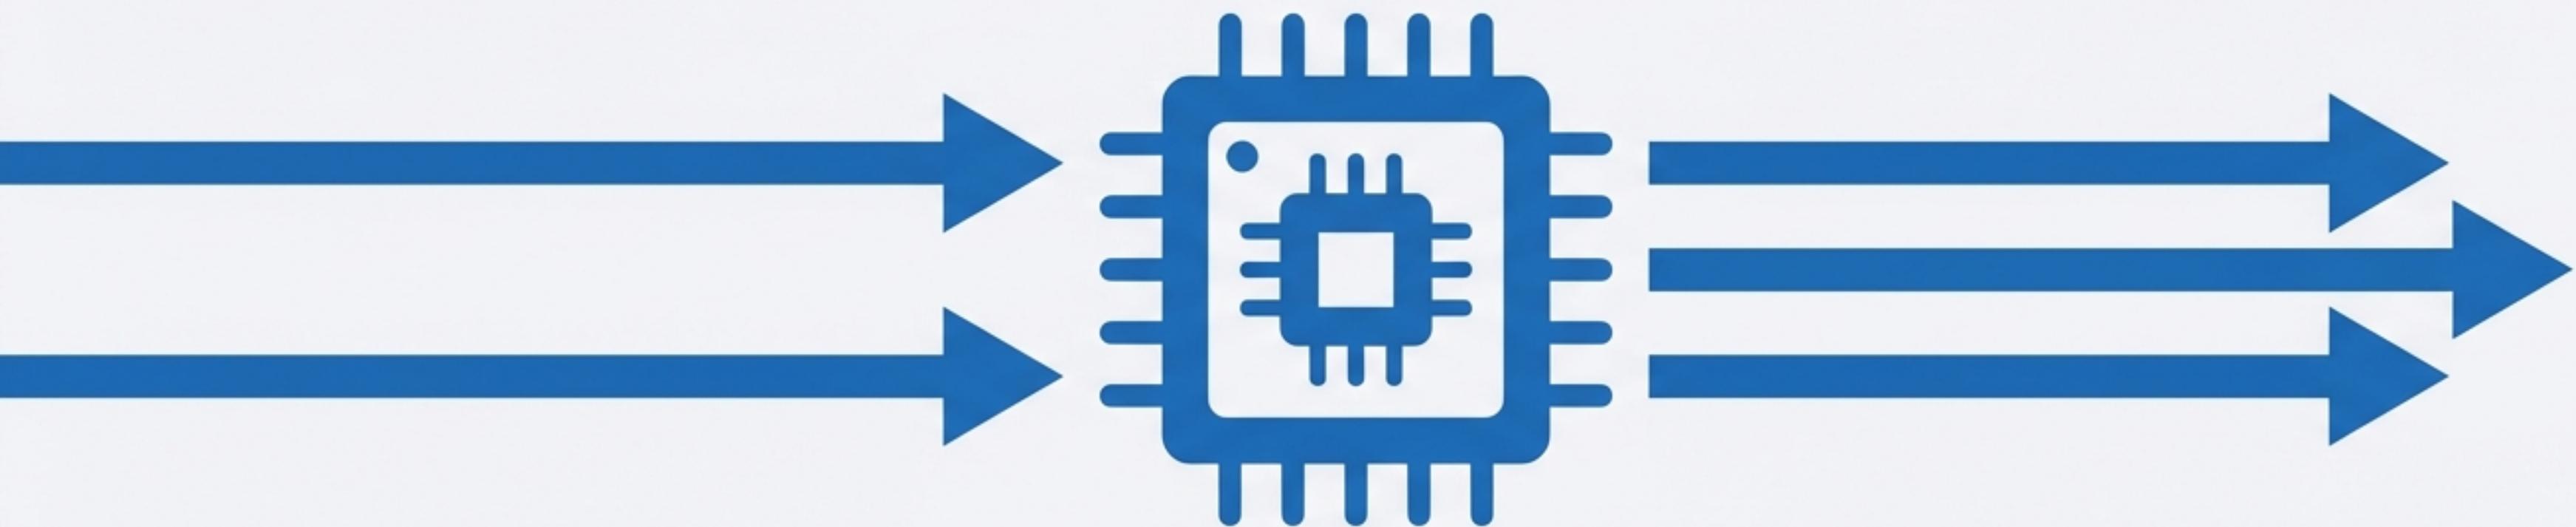


A Database Is a Program, Not Just a Collection of Tables.

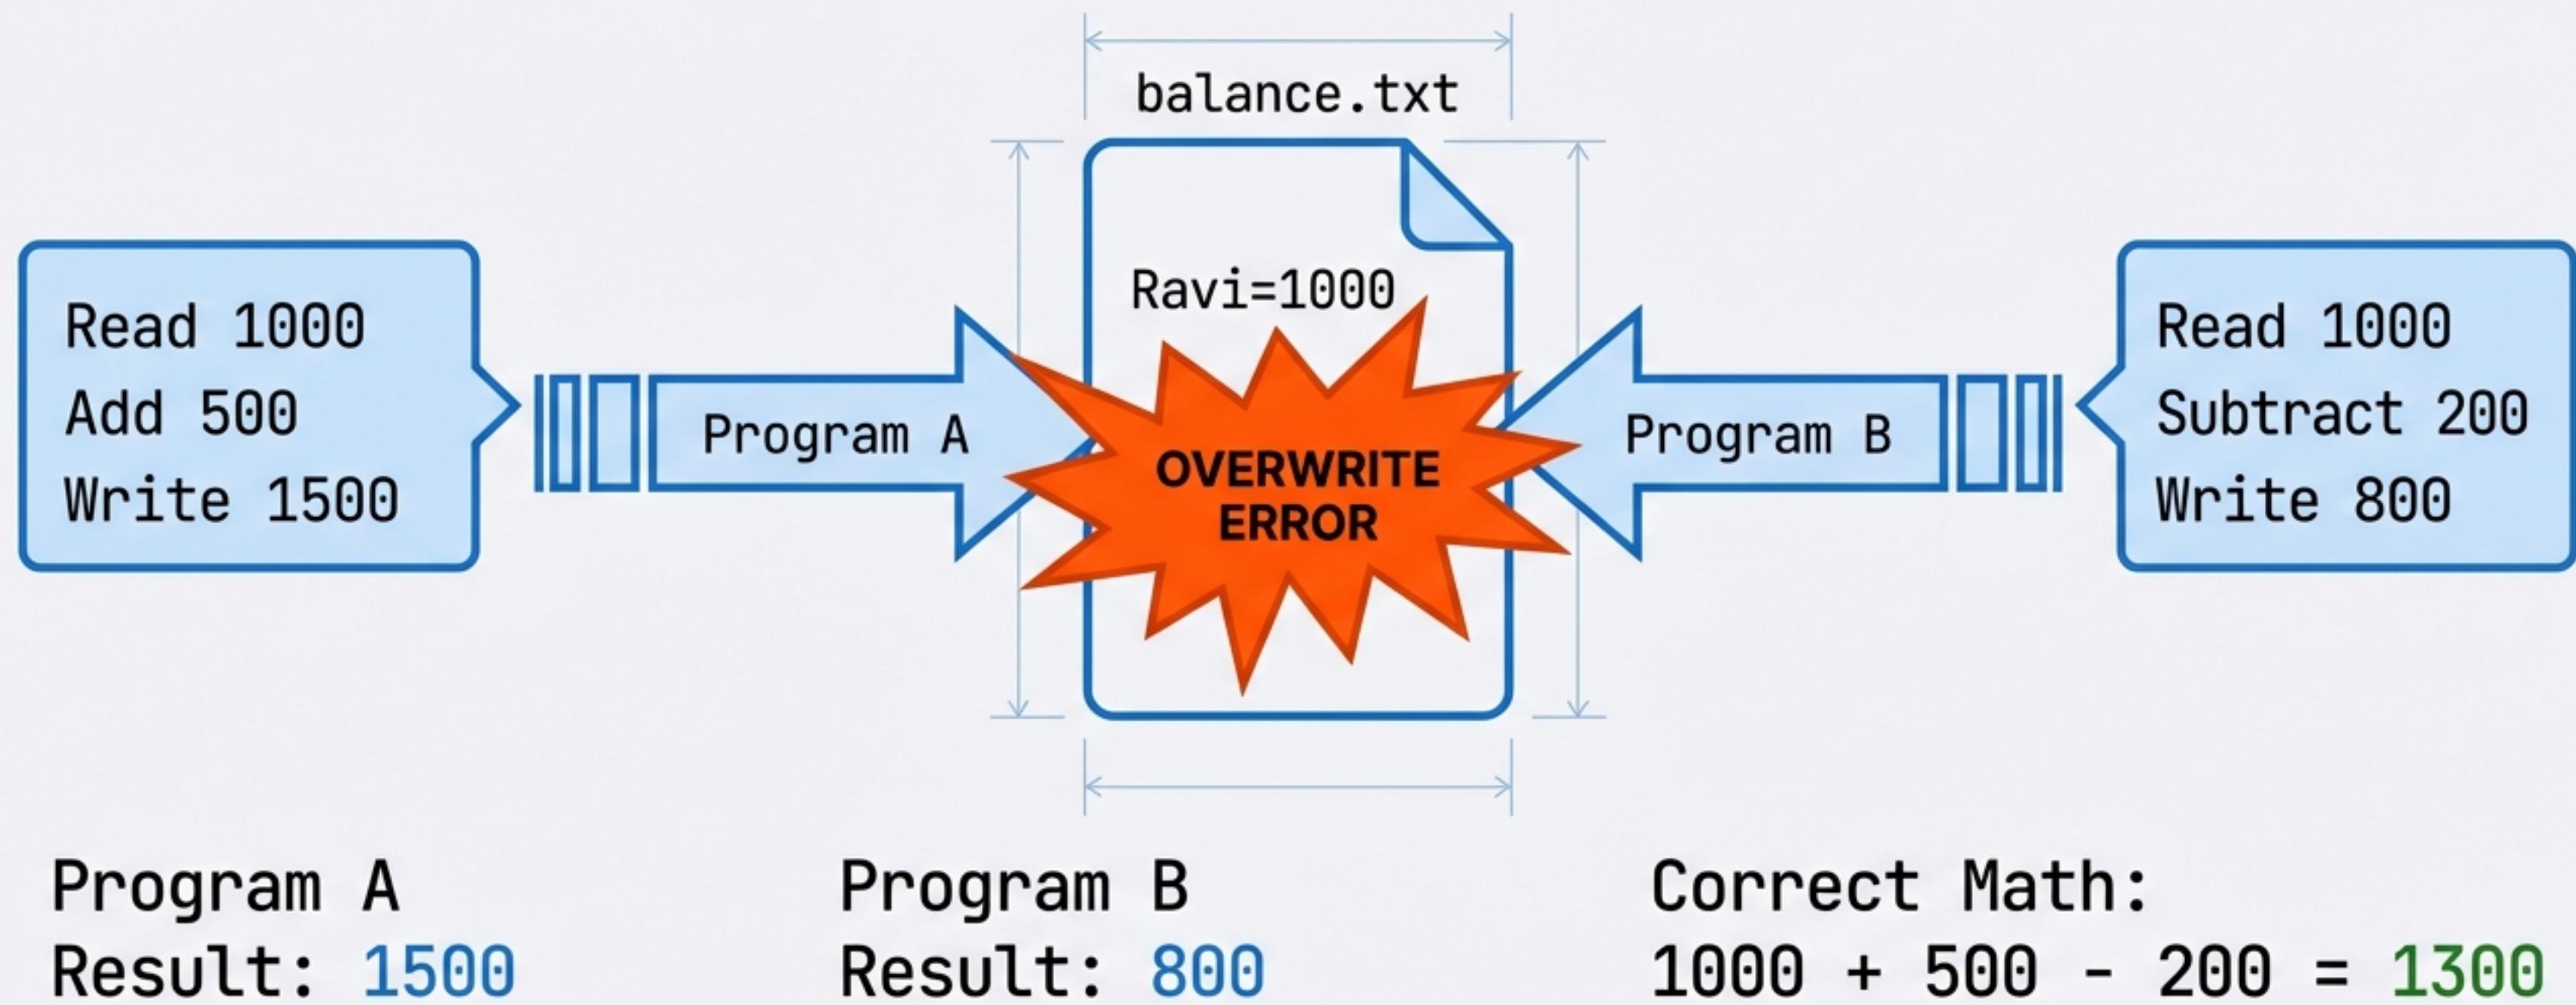


The Core Definition

A database is a running process whose job is to safely manage shared data. It is not just tables, SQL, or indexes—those are tools. The database is the active ‘Gatekeeper’ program.

ETYMOLOGY: DATA + BASE // The concept of one central base where all shared data is stored and managed.

The Overwrite Problem: Why Filesystems Fail at Sharing



Program A
Result: **1500**

Program B
Result: **800**

Correct Math:
 $1000 + 500 - 200 = \text{1300}$

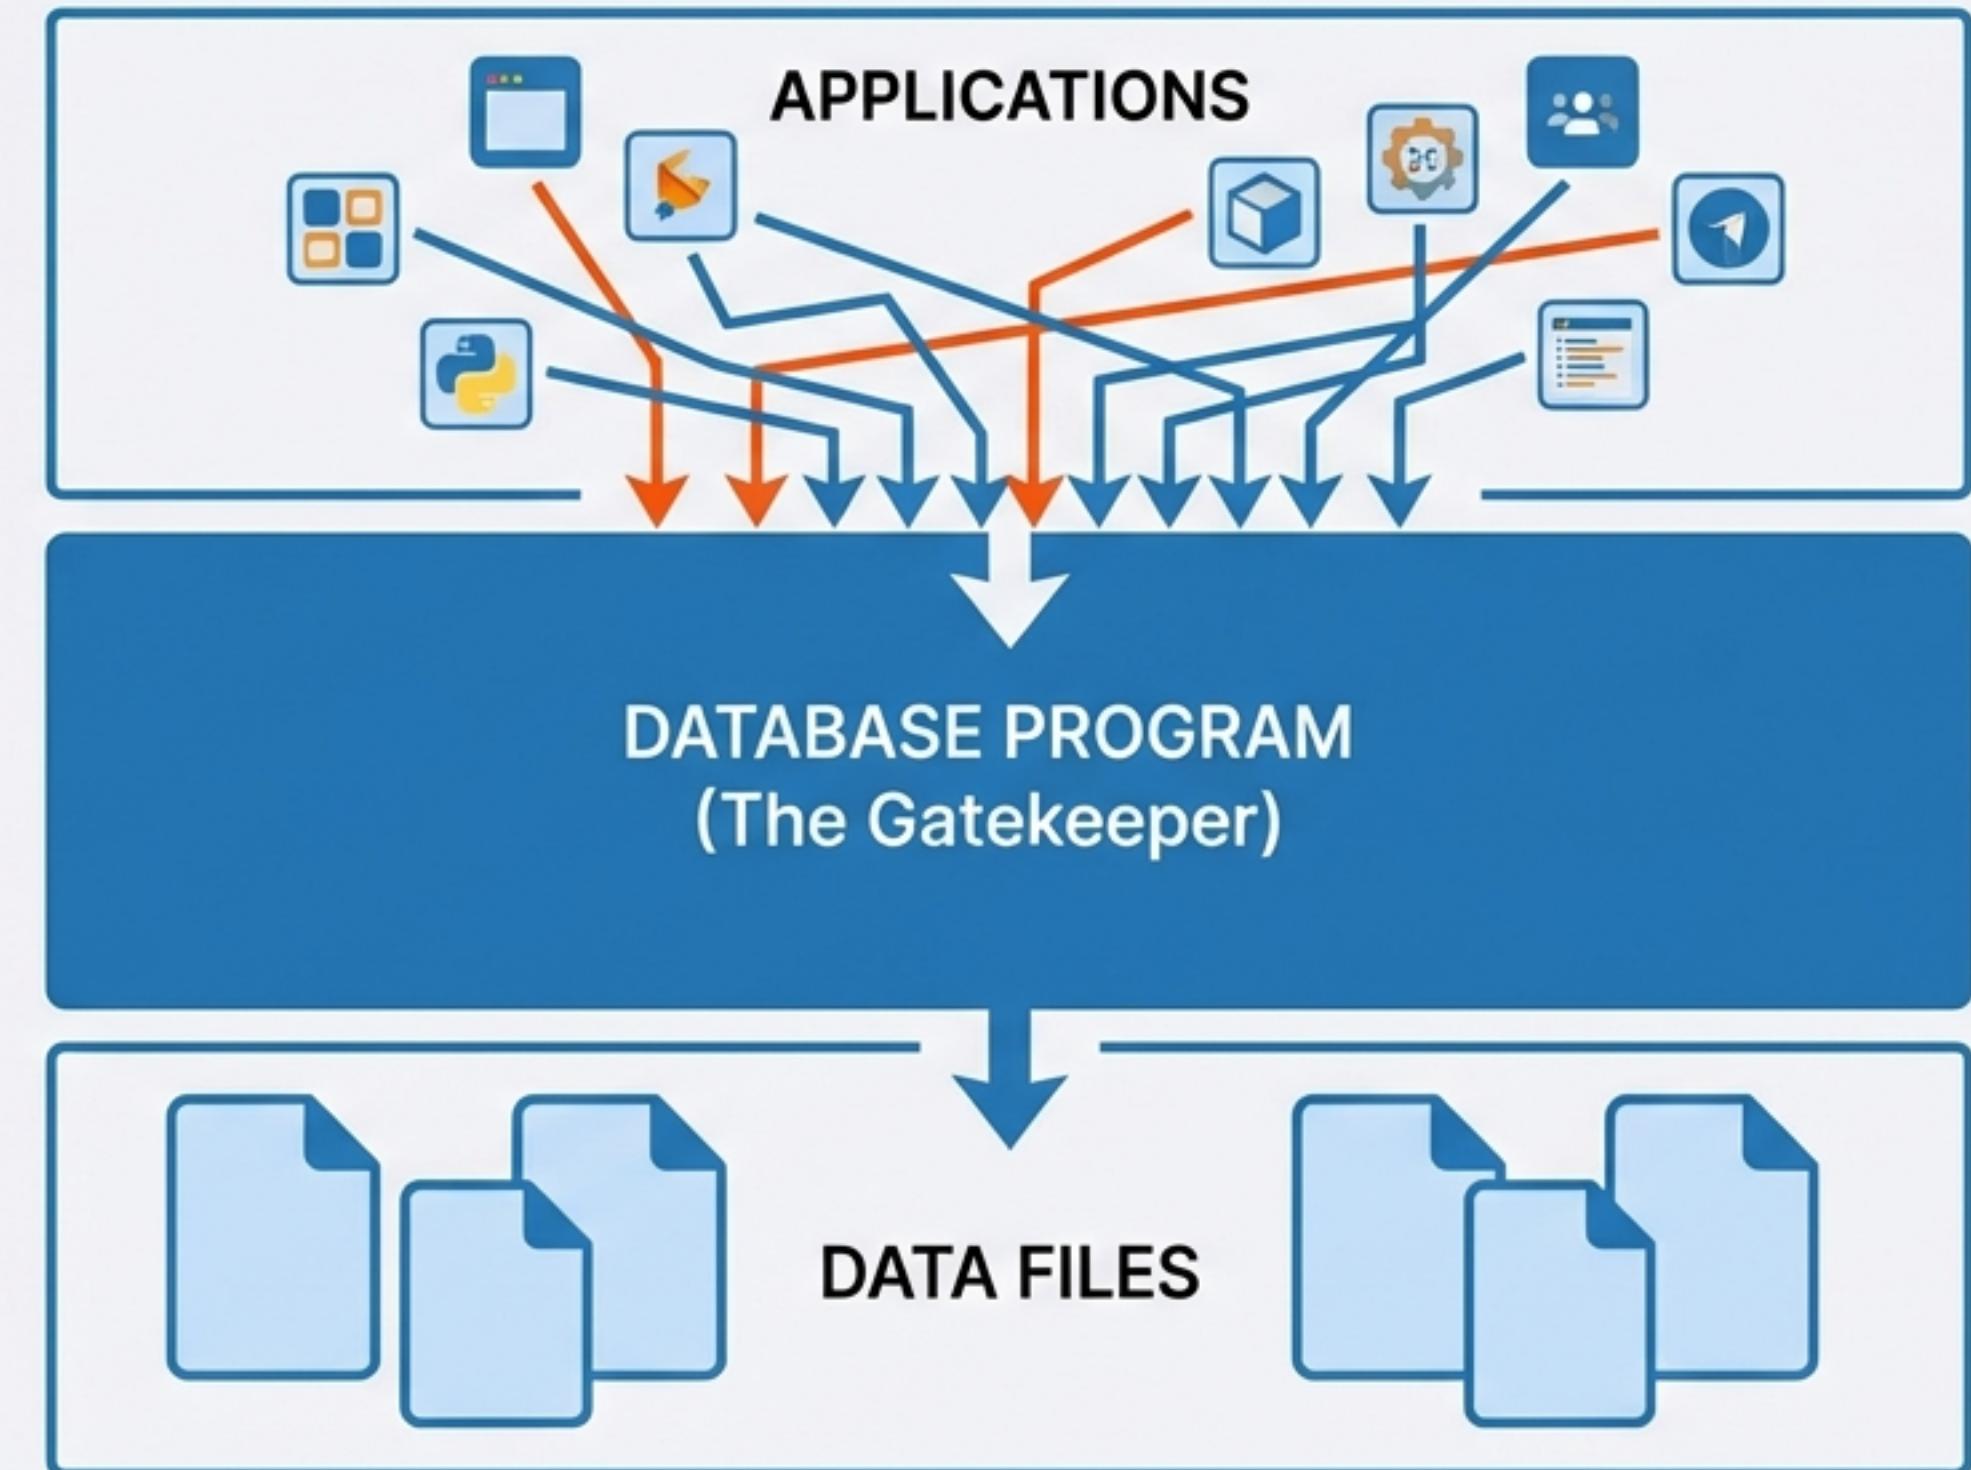
Without a coordinator, the filesystem allows users to unknowingly destroy each other's work.

The Architecture Shift: Introducing the Gatekeeper

The New Rule: No application touches files directly.

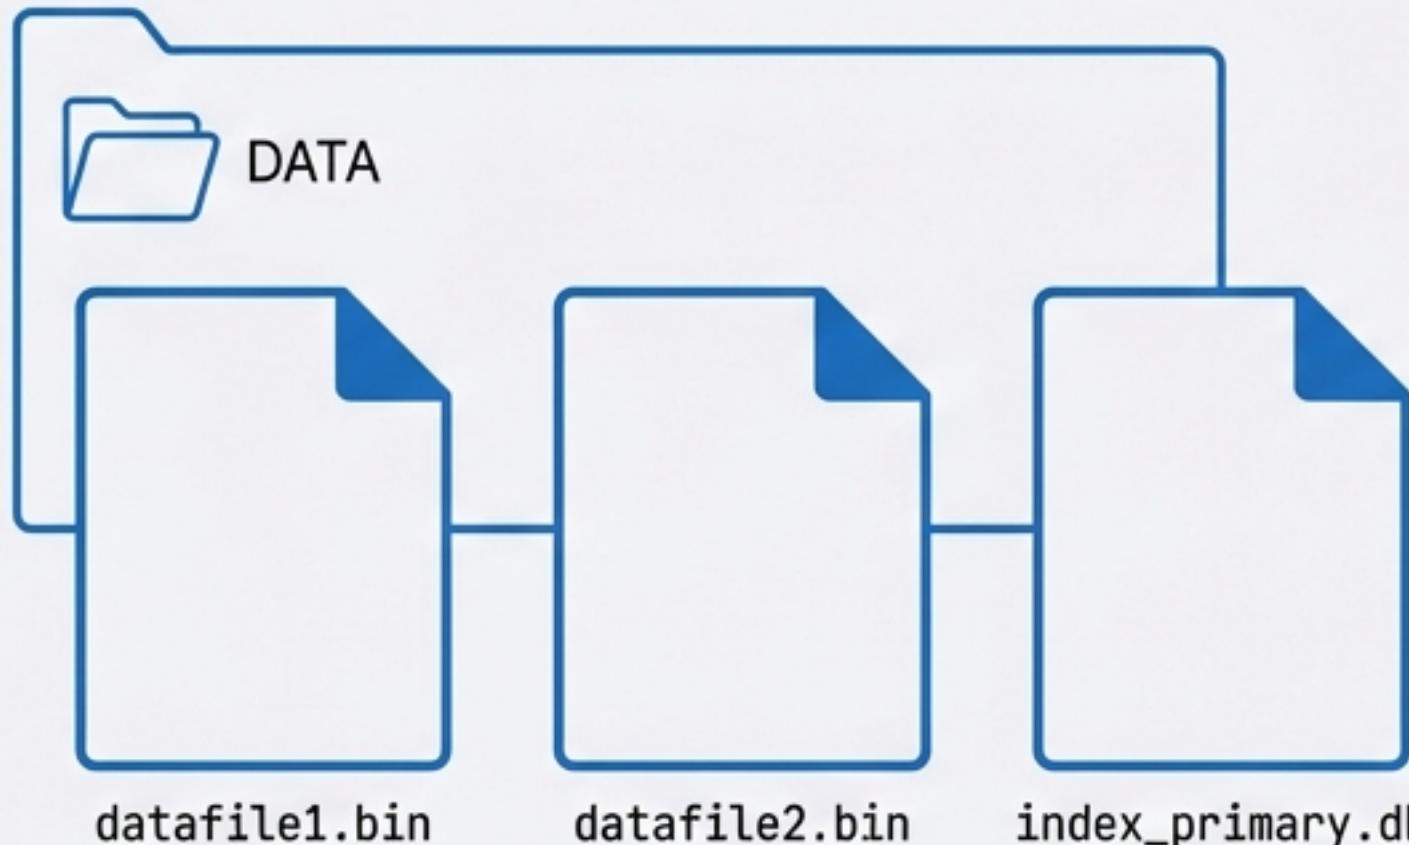
The Mechanism: The database process serializes requests. It forces simultaneous writes into a single line.

// Outcome: Only ONE program controls writes to shared data.



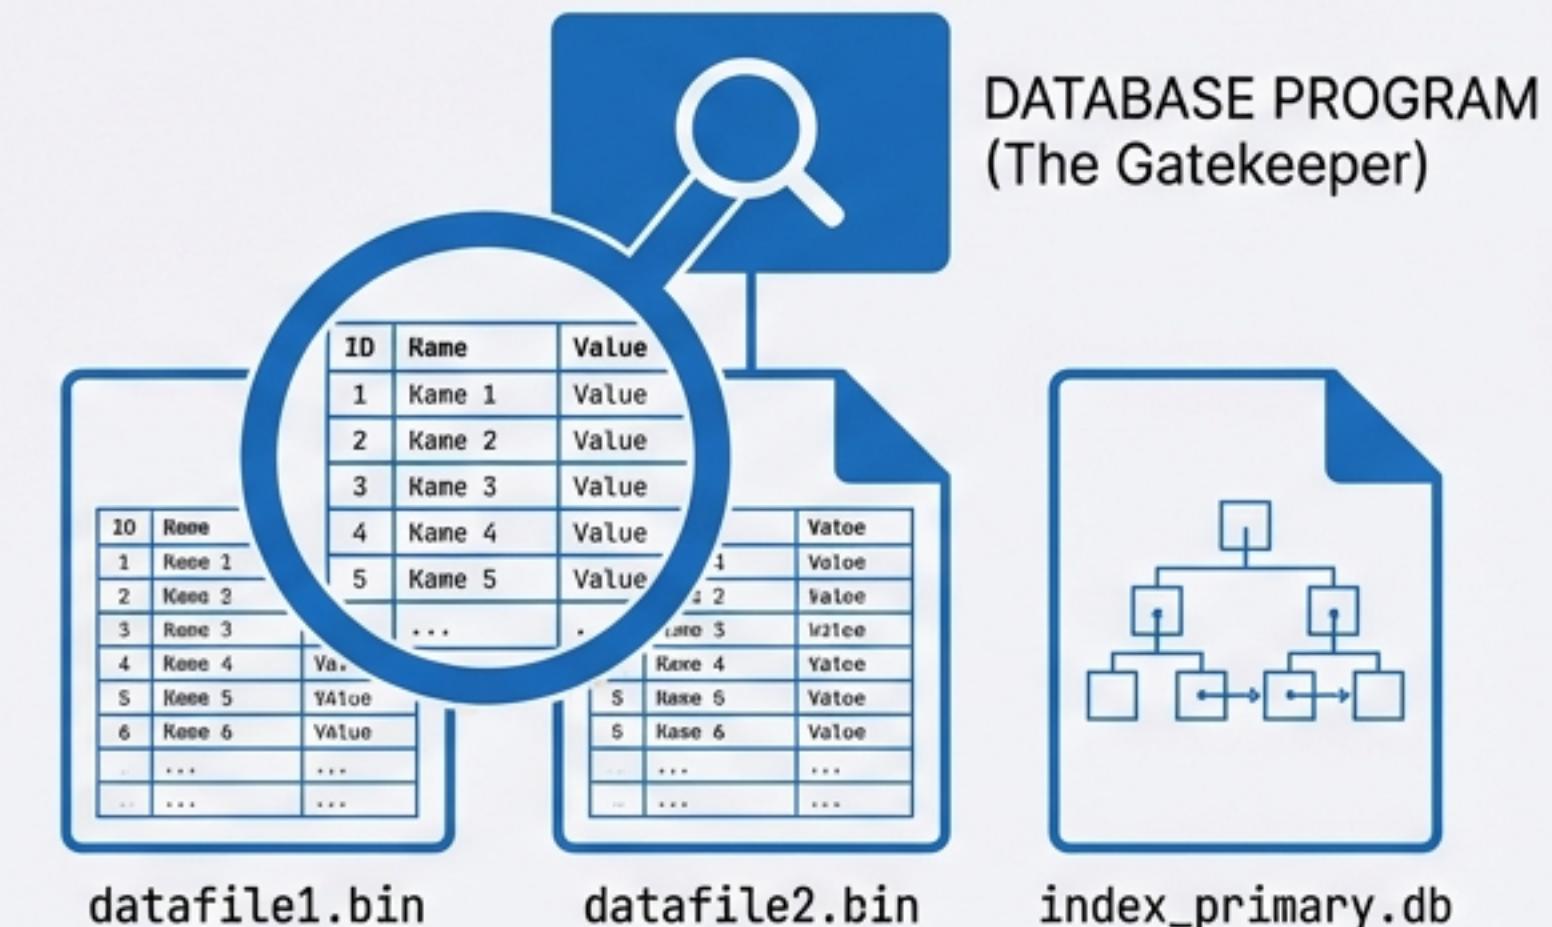
The Database Still Uses Files Underneath

THE FILESYSTEM VIEW



Stores bytes. Names files.
Does **not** know meaning.

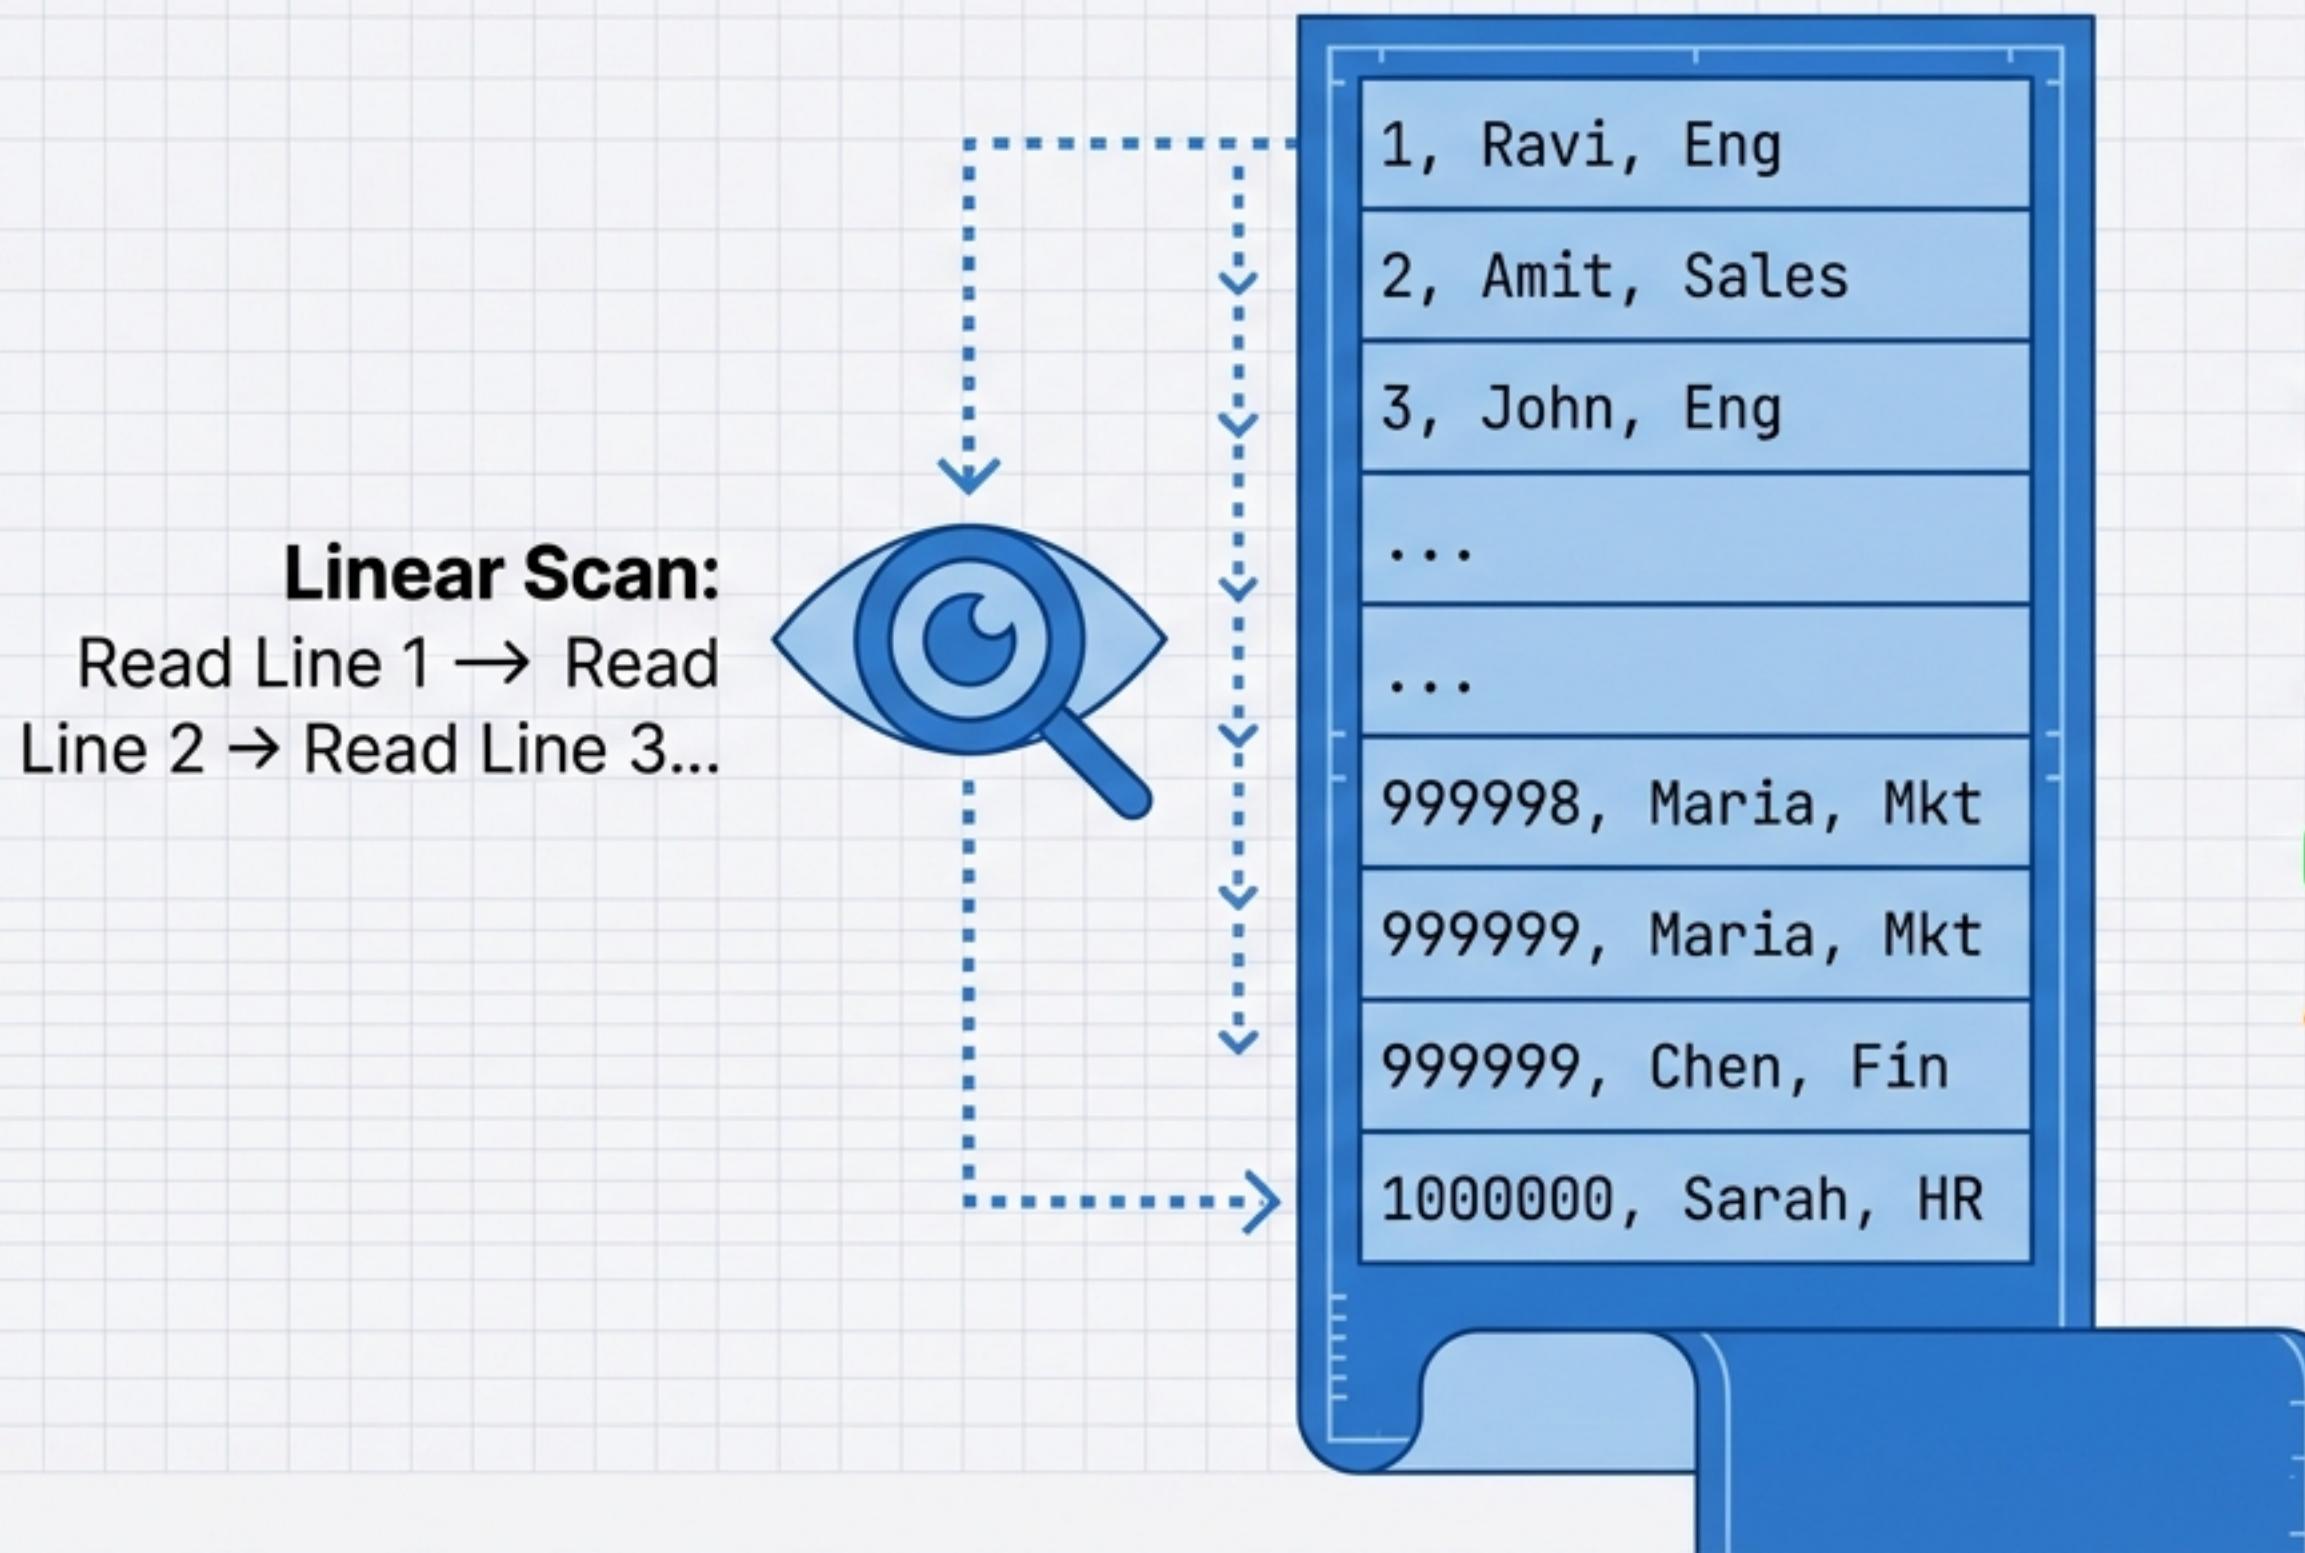
THE DATABASE VIEW



Understands structure.
Controls access. Enforces rules.

The database doesn't remove files; it wraps them in logic.

Safety Was Solved, but Search Was Slow



The Problem:

To find one record, the computer must read every single line.

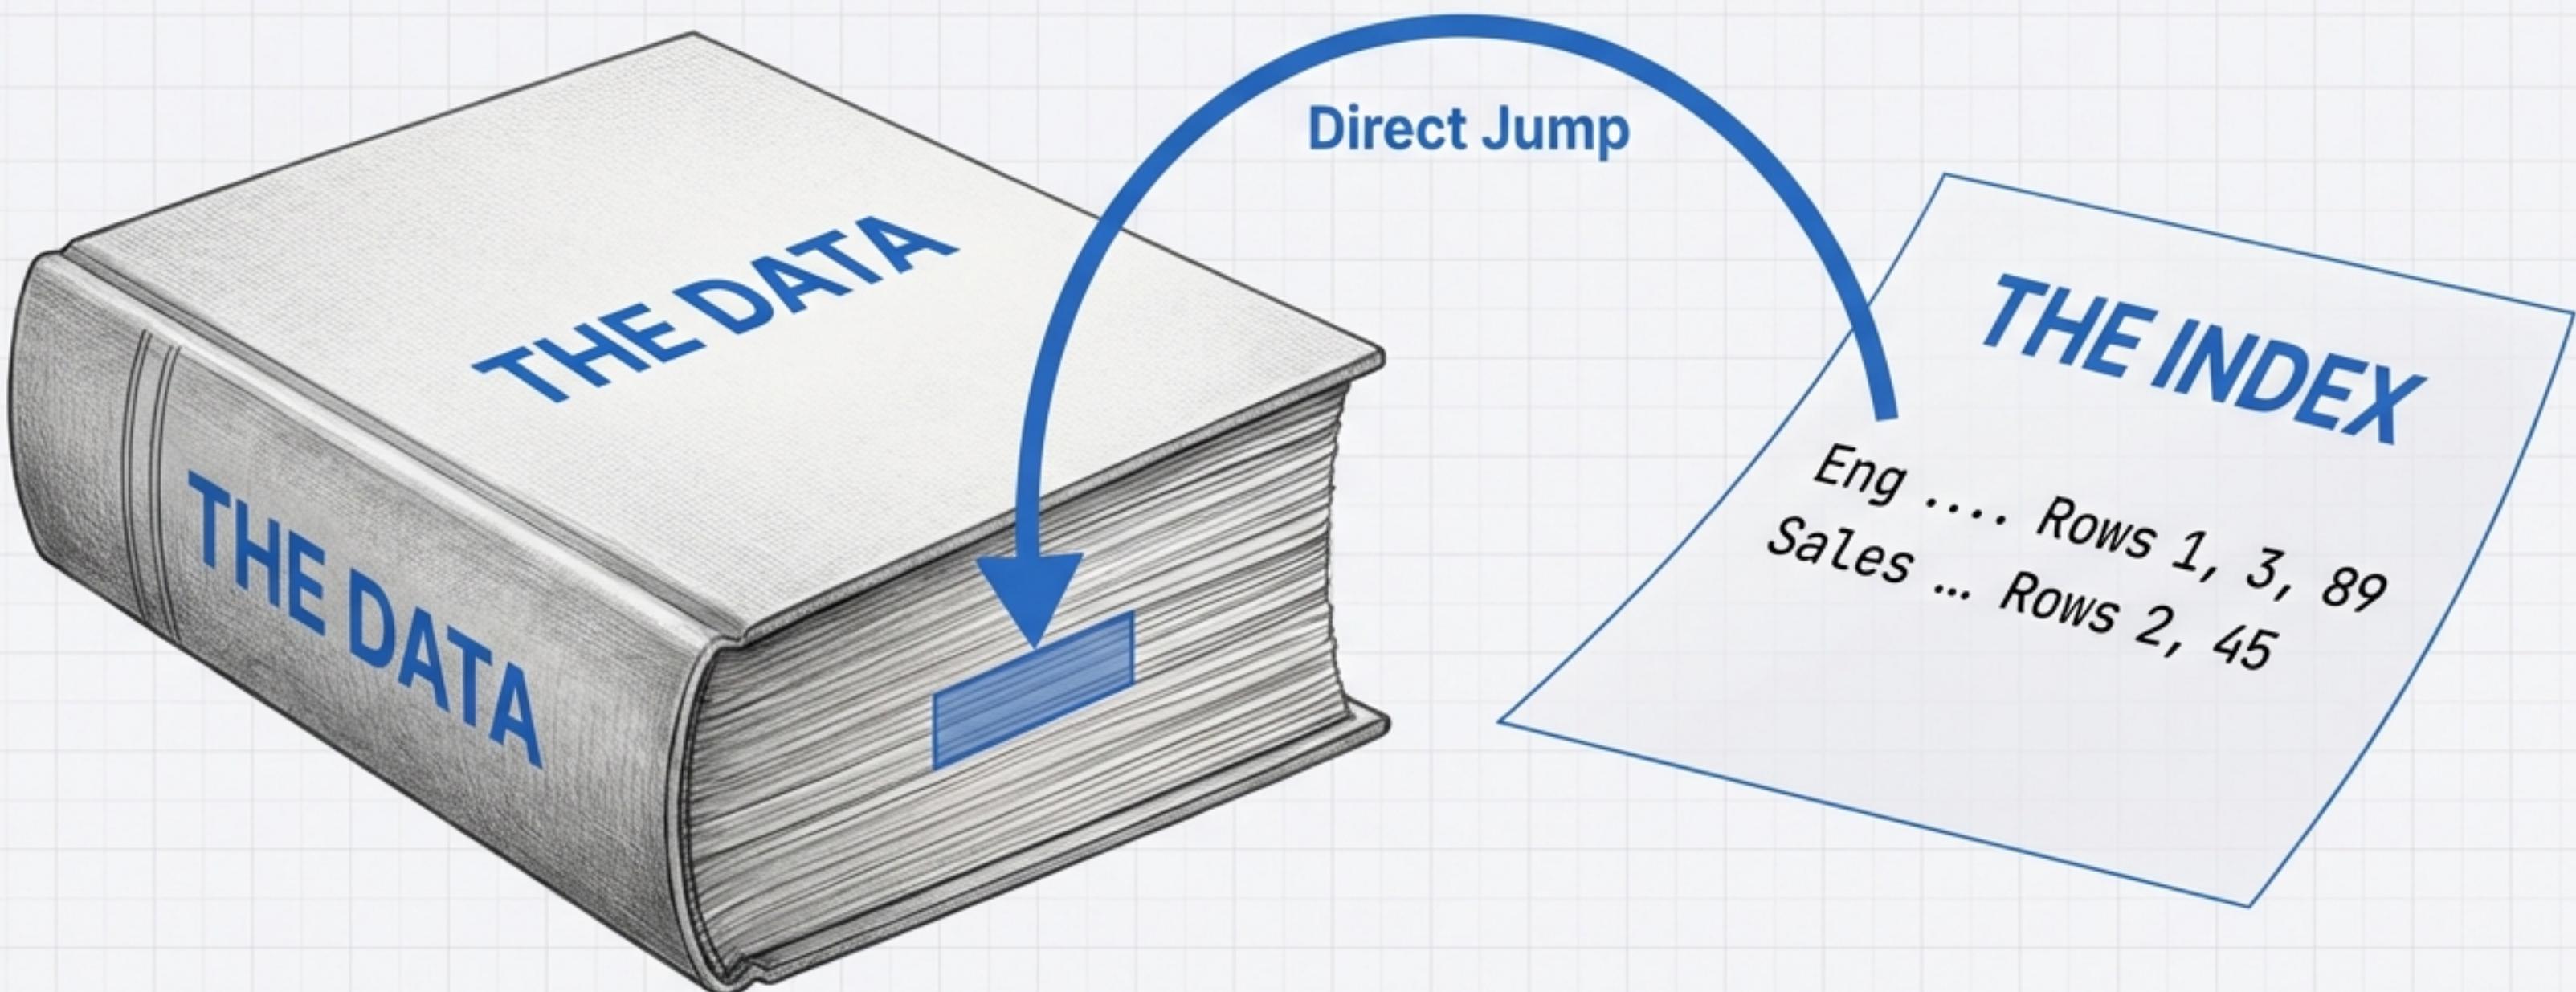
The Scale:

Linear Scan:
Read Line 1 → Read Line 2 → Read Line 3...

✓ 10 Rows = Fast

⚠ 10 Million Rows = Disaster

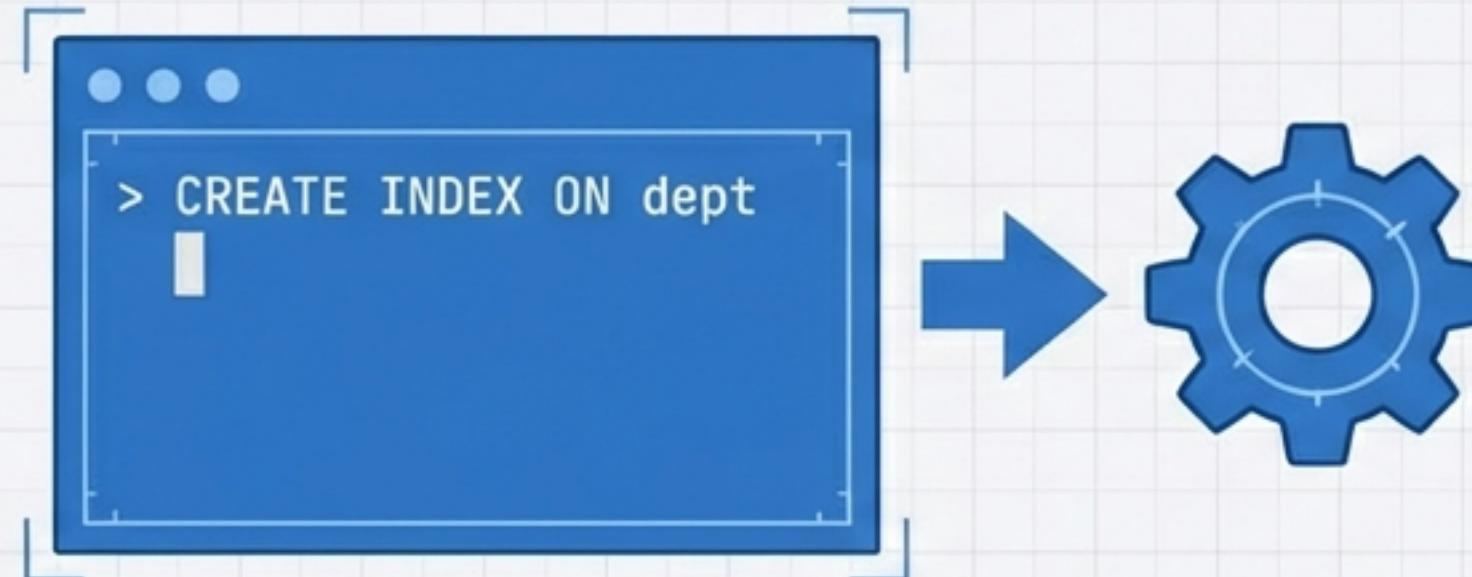
The Solution: An Index Is a Shortcut, Not the Data



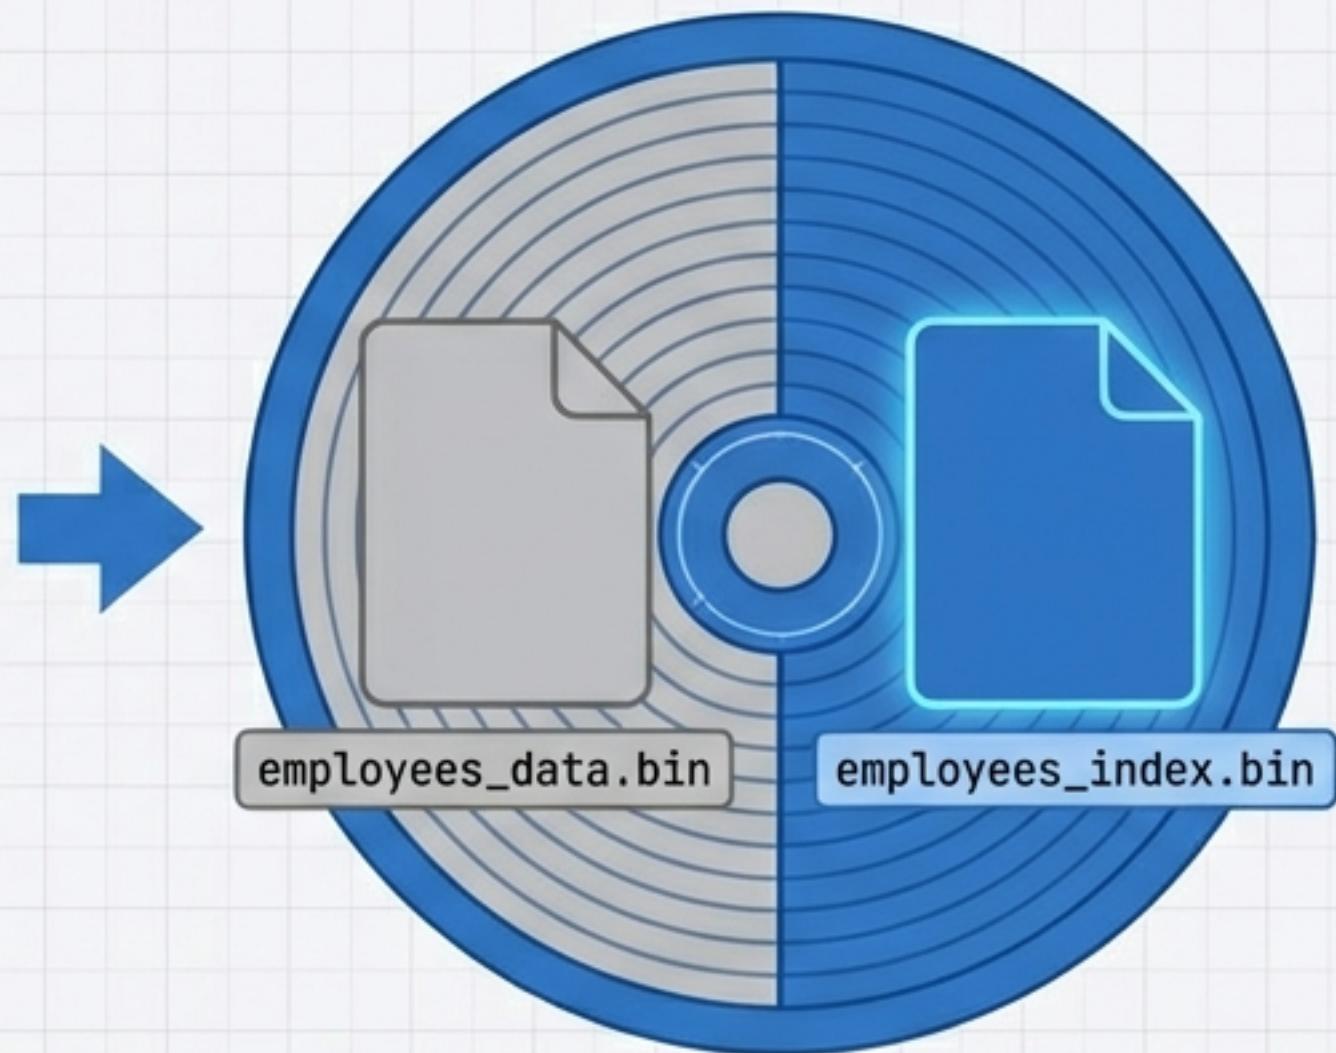
Concept: You don't read the whole book; you check the index and jump directly to the page.

Implementation: The database maintains a secondary structure that maps "Value" → "Row Location".

How the Database Builds an Index



Database scans
table ONCE.
Builds mapping:
Value → Row ID.



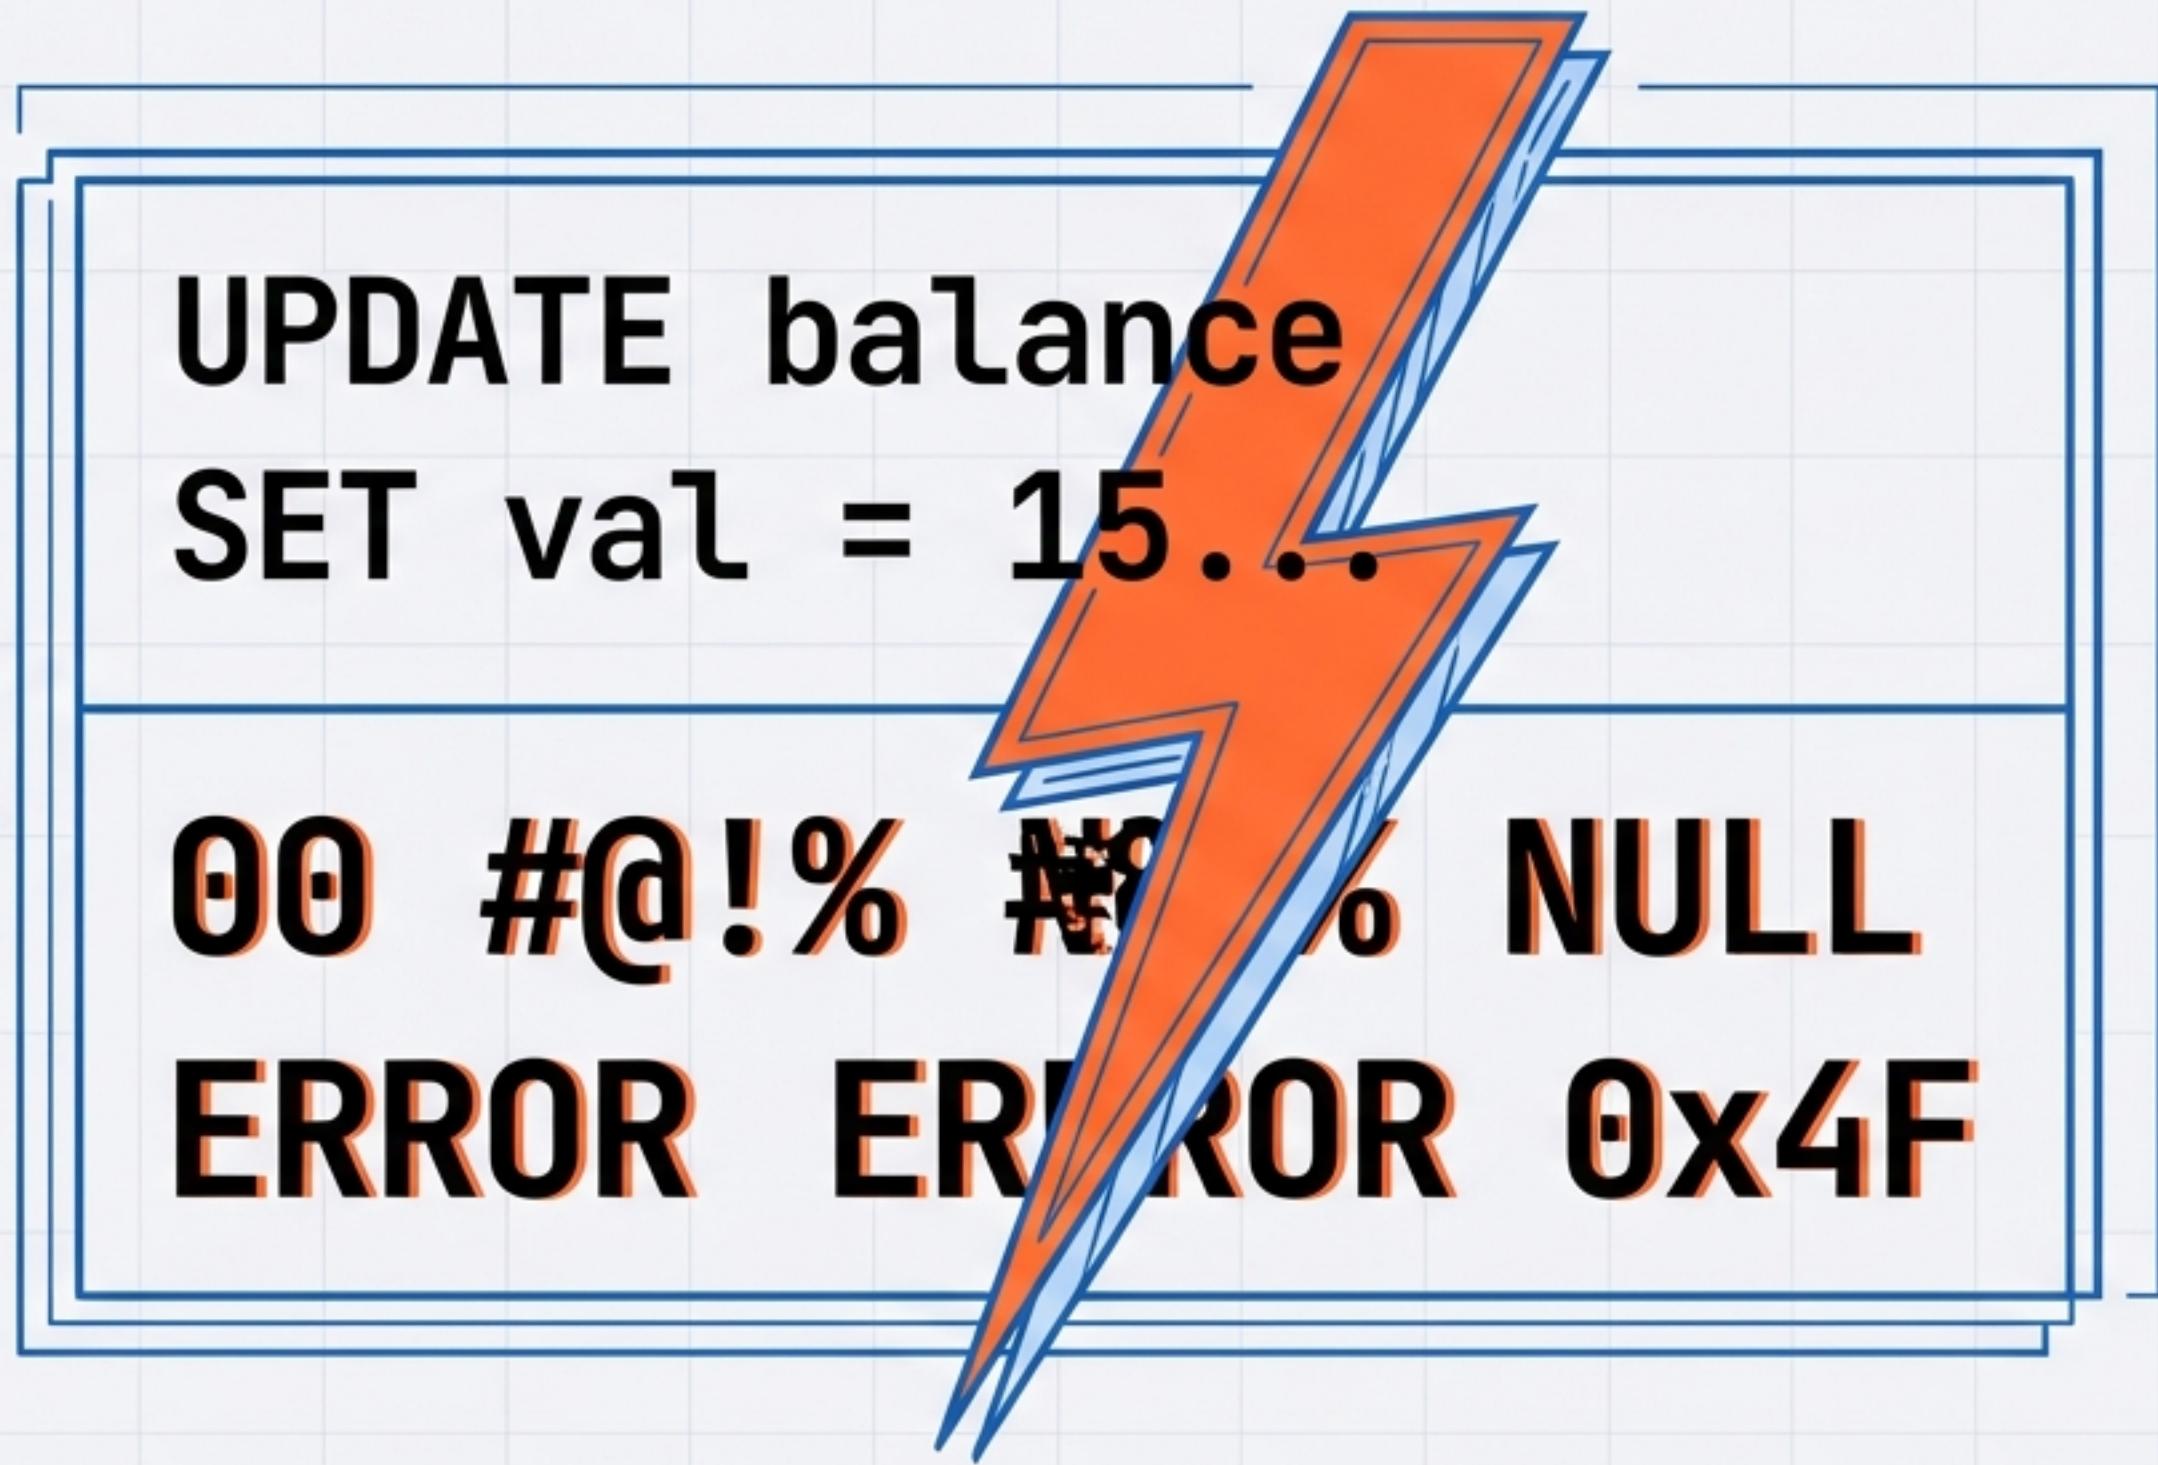
Crucial Insight: The Index is physically stored as a SEPARATE file that the database must update and maintain.

The Crash Problem: When Writes Fail Halfway

The Scenario: Updating Balance from 1000 to 1500.

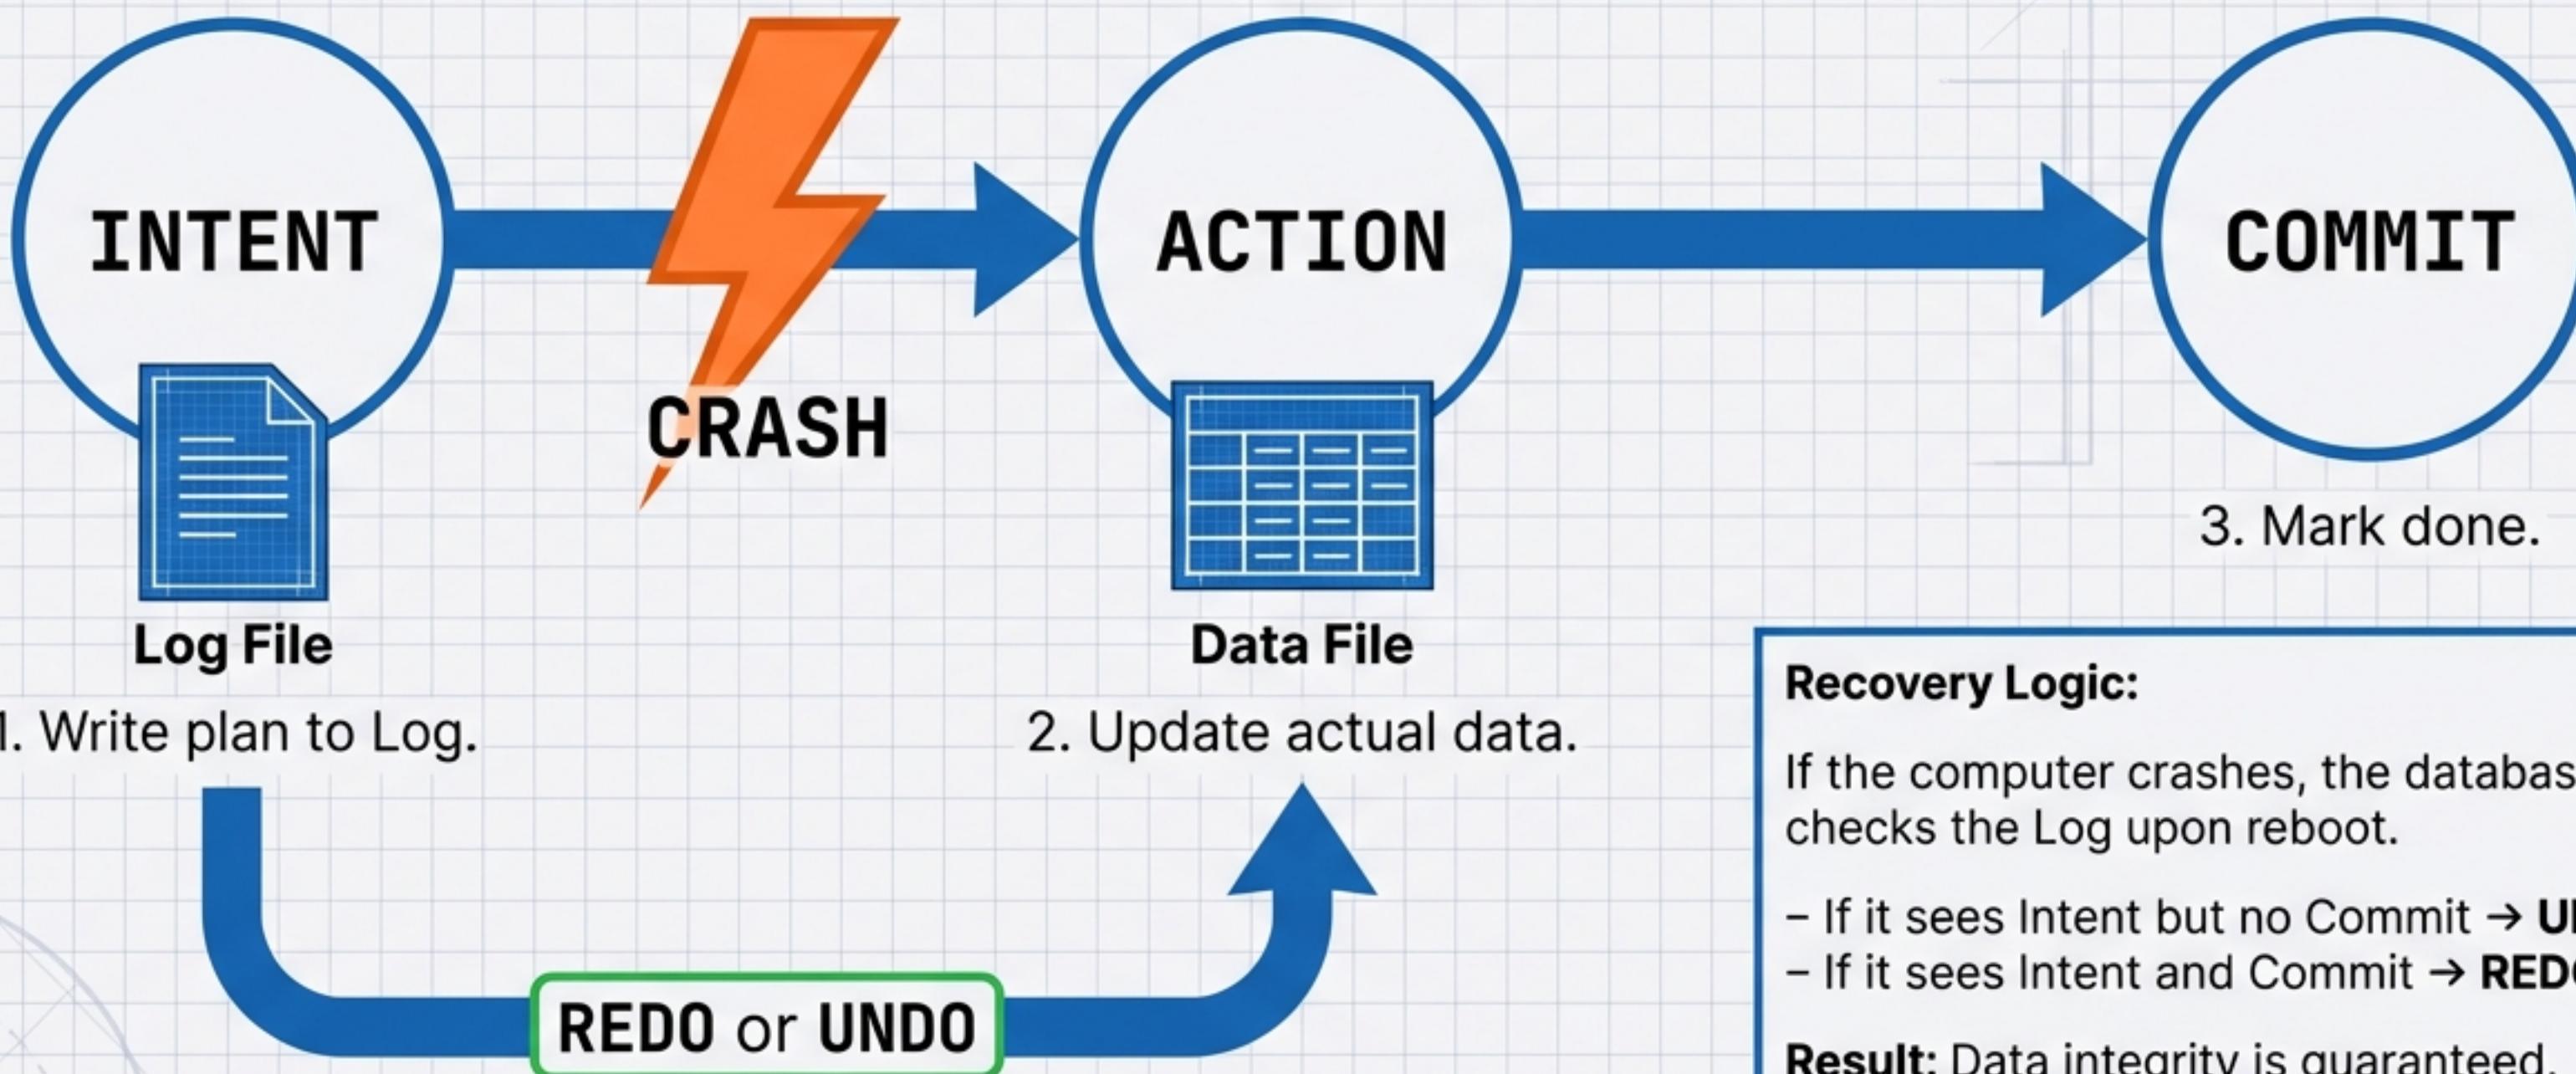
The Disaster: Power fails while writing the bytes.

The Result: The file contains half-written garbage. The filesystem cannot recover this because it doesn't know what you were trying to do.



The Solution: Write the Intent Before the Action

The Write-Ahead Log (WAL)



Recovery Logic:

If the computer crashes, the database checks the Log upon reboot.

- If it sees Intent but no Commit → UNDO.
- If it sees Intent and Commit → REDO.

Result: Data integrity is guaranteed.

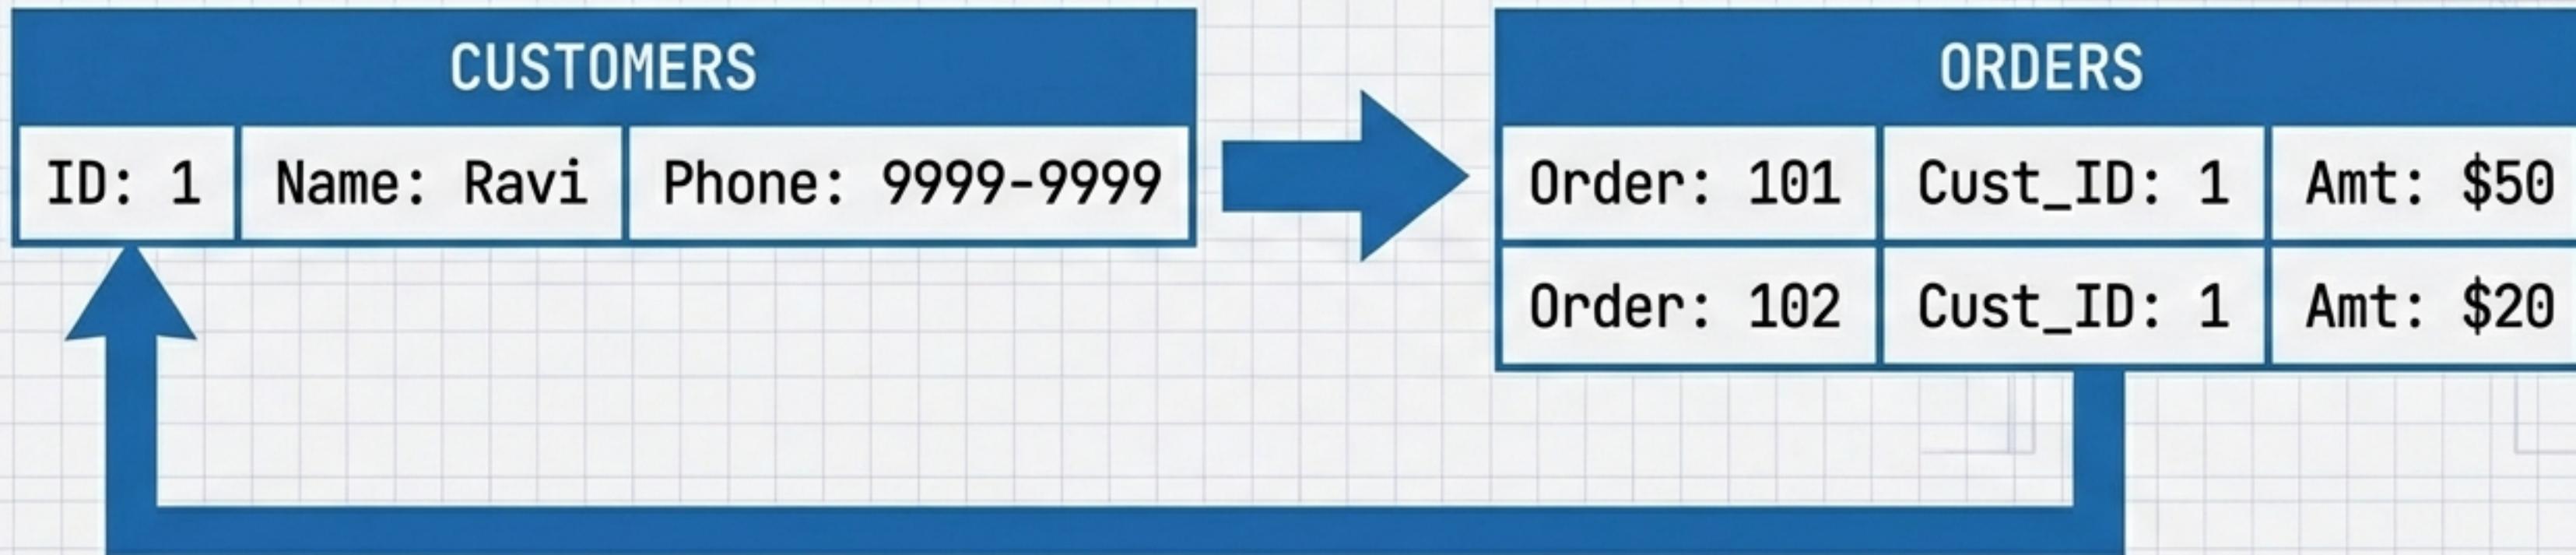
The Consistency Problem: The Nightmare of Flat Files

ORDERS_FLAT_FILE		
ID:	Customer:	Phone:
1	Ravi	9999-9999
2	Ravi	9999-9999
3	Ravi	5555-0000

Error: Ravi's phone number was updated here, but not in Rows 1 and 2.

Data Duplication leads to Data Corruption. In a flat file, information is repeated. If you forget to update one instance, the database contradicts itself. Which phone number is real?

The Relational Model: Splitting Data by Meaning

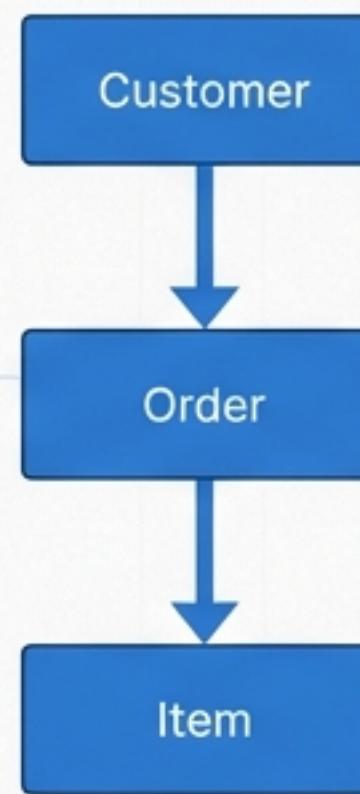


1. Split data into logical themes (Customers vs Orders).
2. Store the customer details exactly ONCE.
3. Reference the customer by ID.

Result: Update the phone number in one place, and it reflects everywhere.

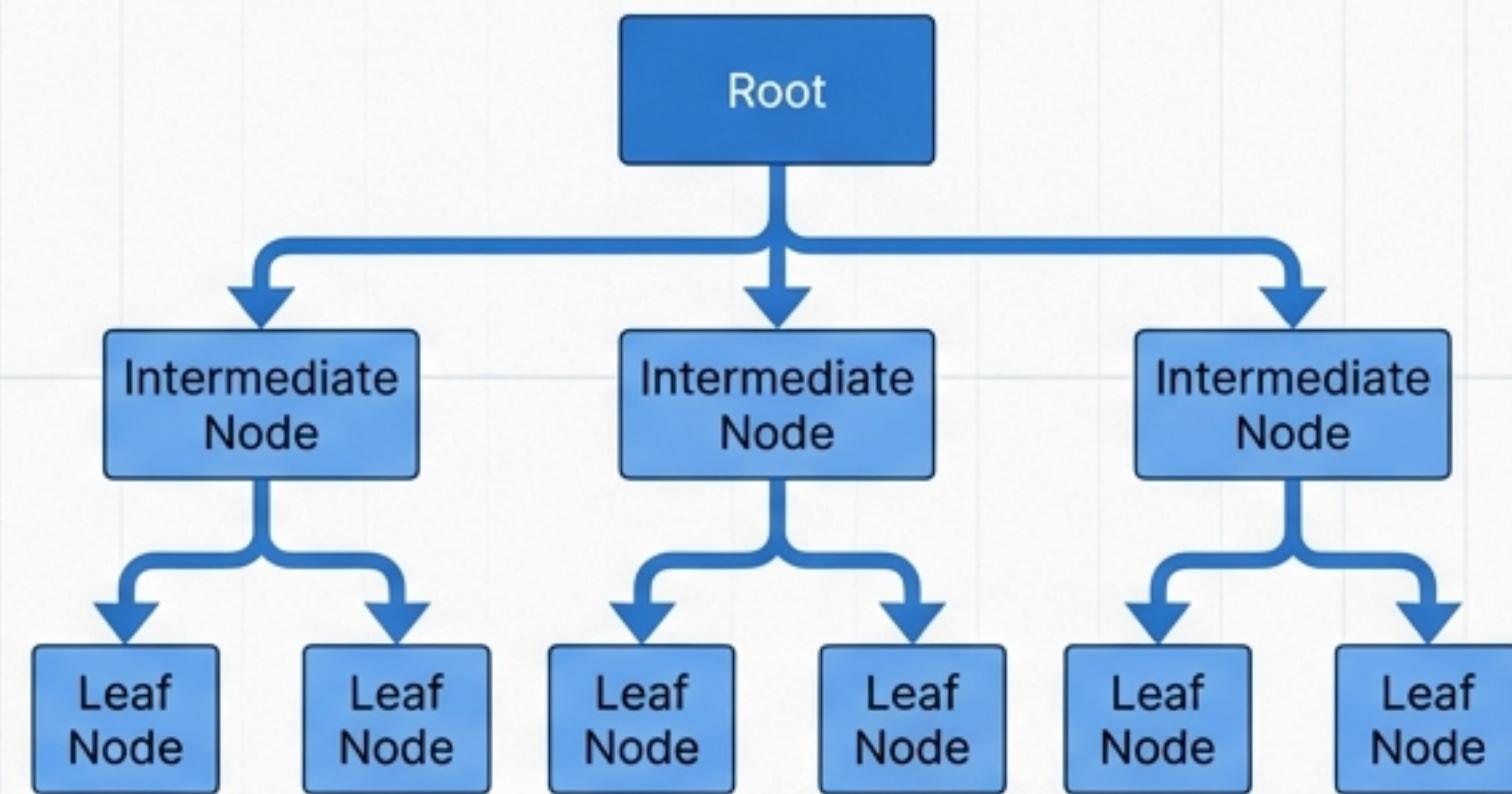
Two Types of “Trees”: Don’t Confuse Them

The Old Hierarchical Model



Rigid. Hard to change.
Pre-dates relational databases.

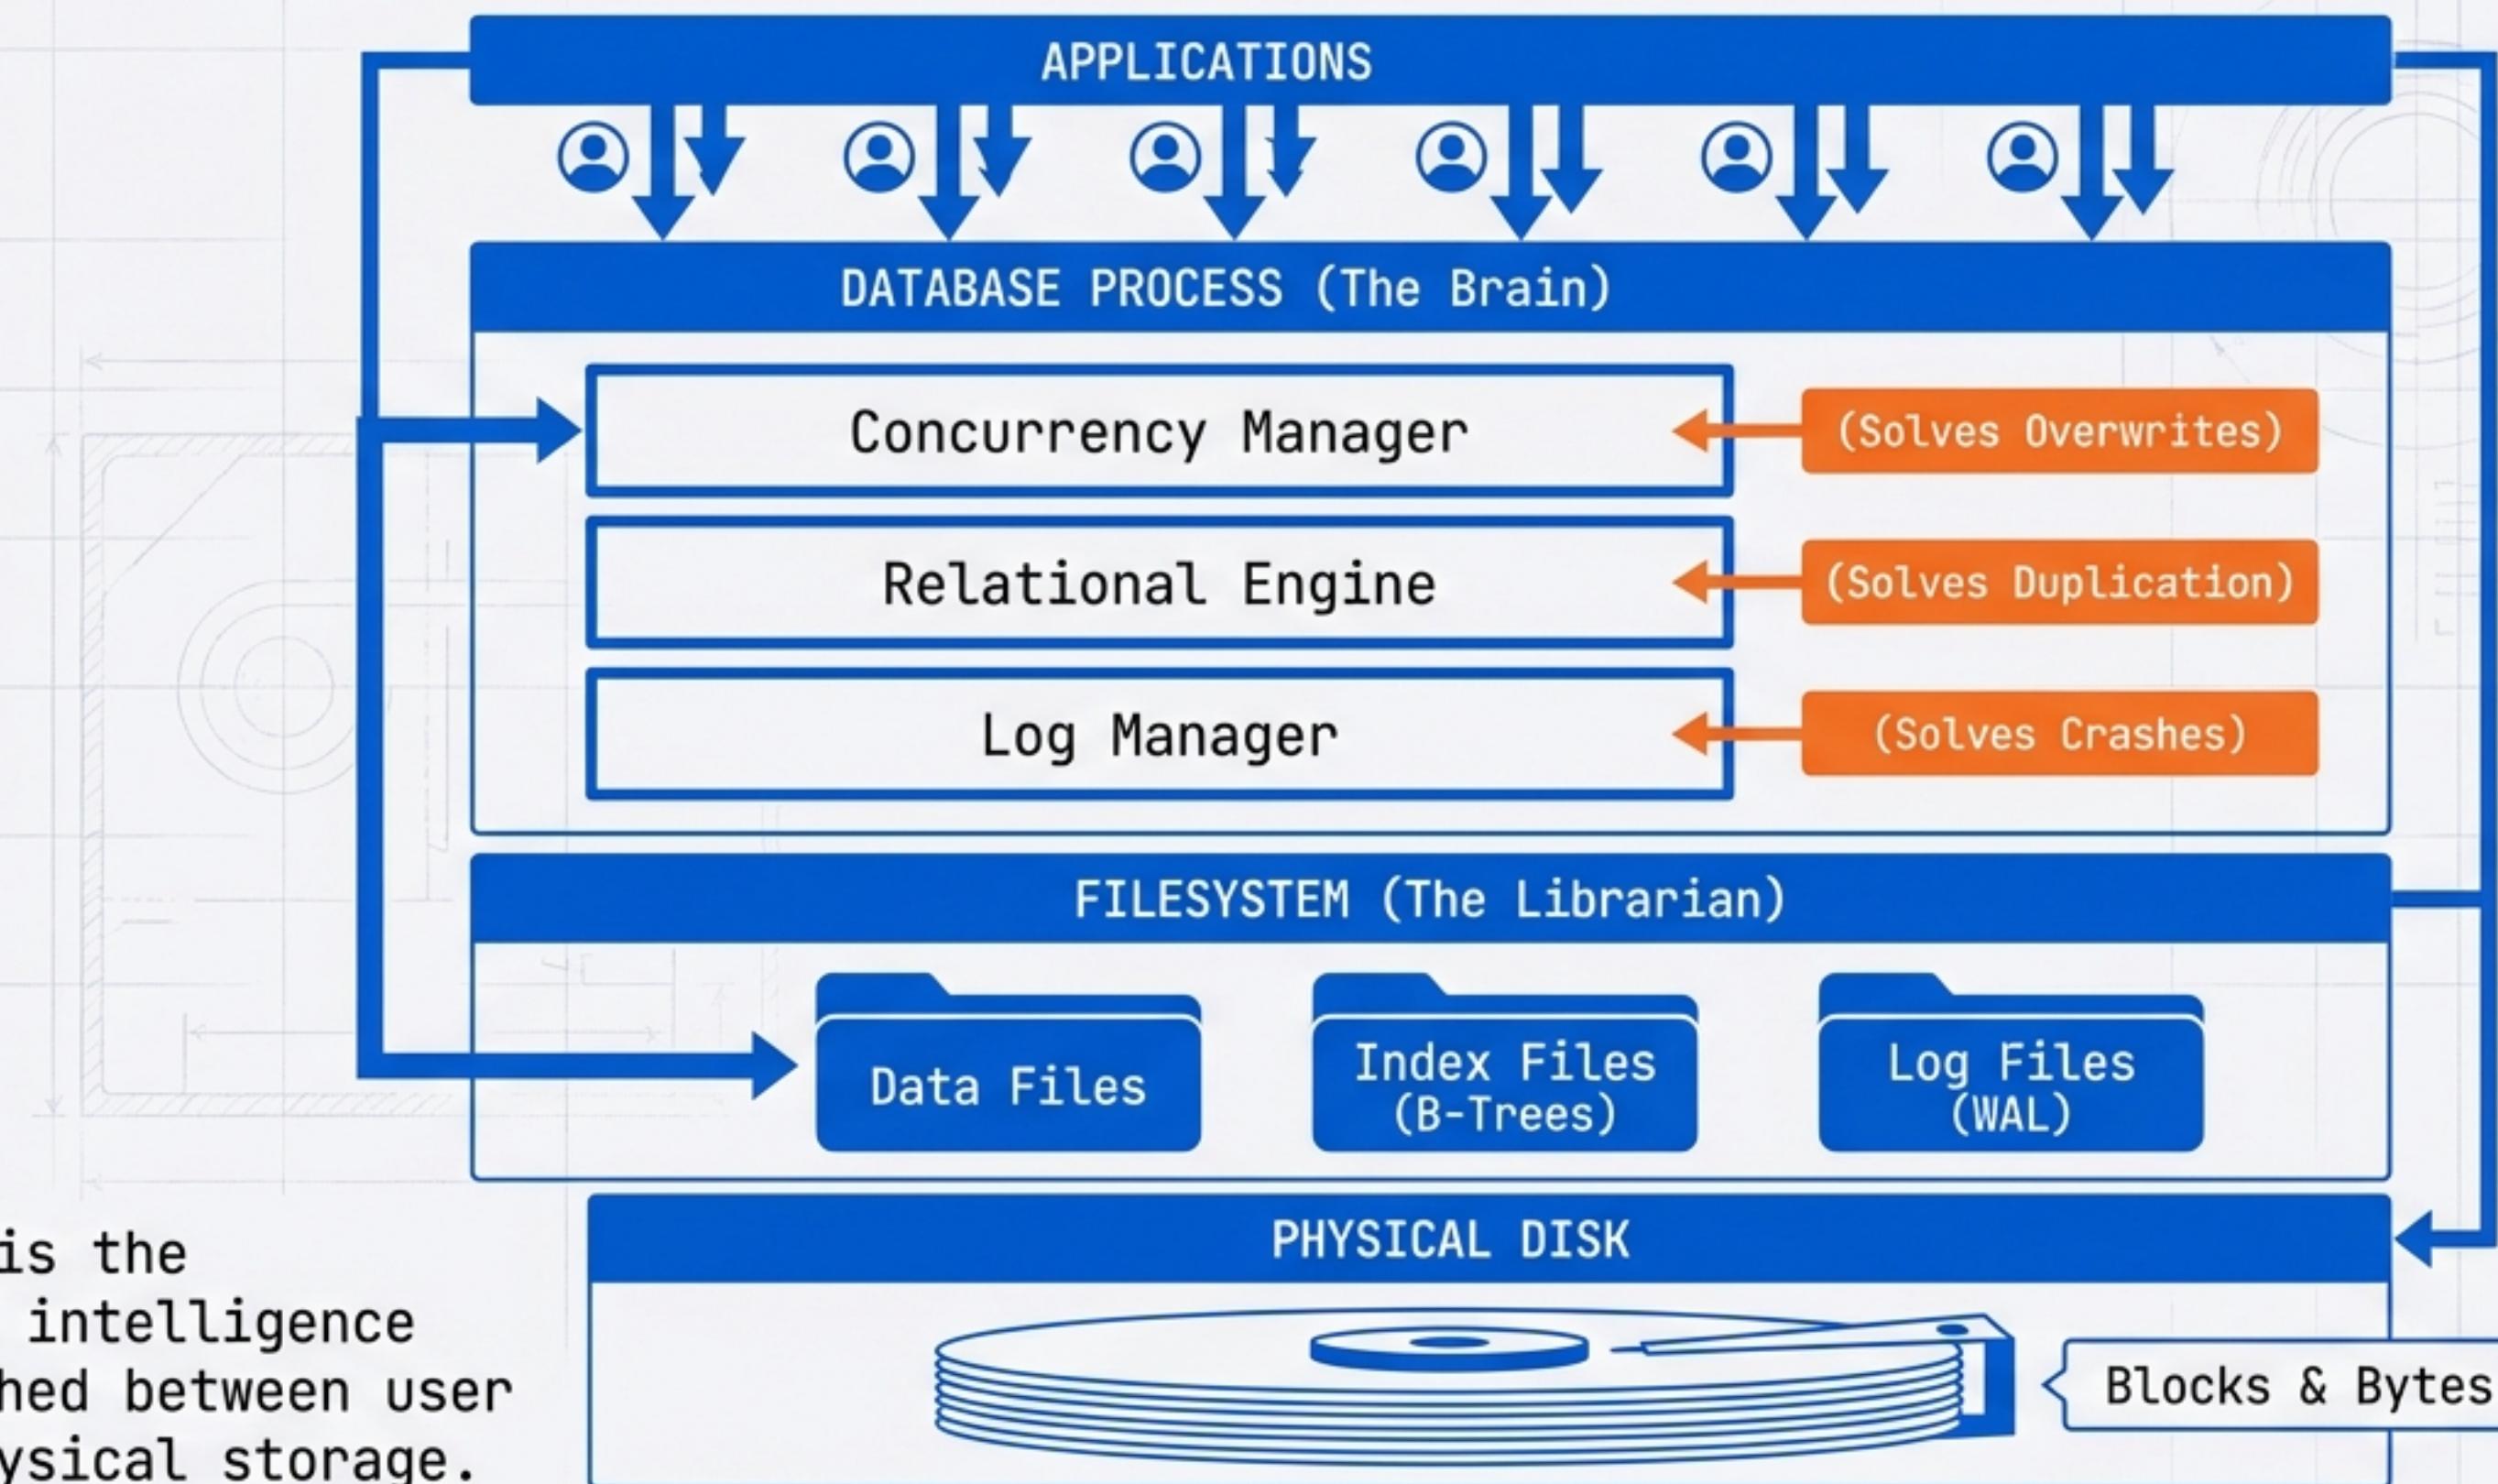
The Modern B-Tree Index



The B-Tree. This is the internal data structure used for Indexes. It is optimized for computer search speed (finding a value in the fewest jumps).

When engineers talk about “Trees” in databases today, they usually mean the B-Tree Index, not the data model.

The Complete Architecture: From App to Disk

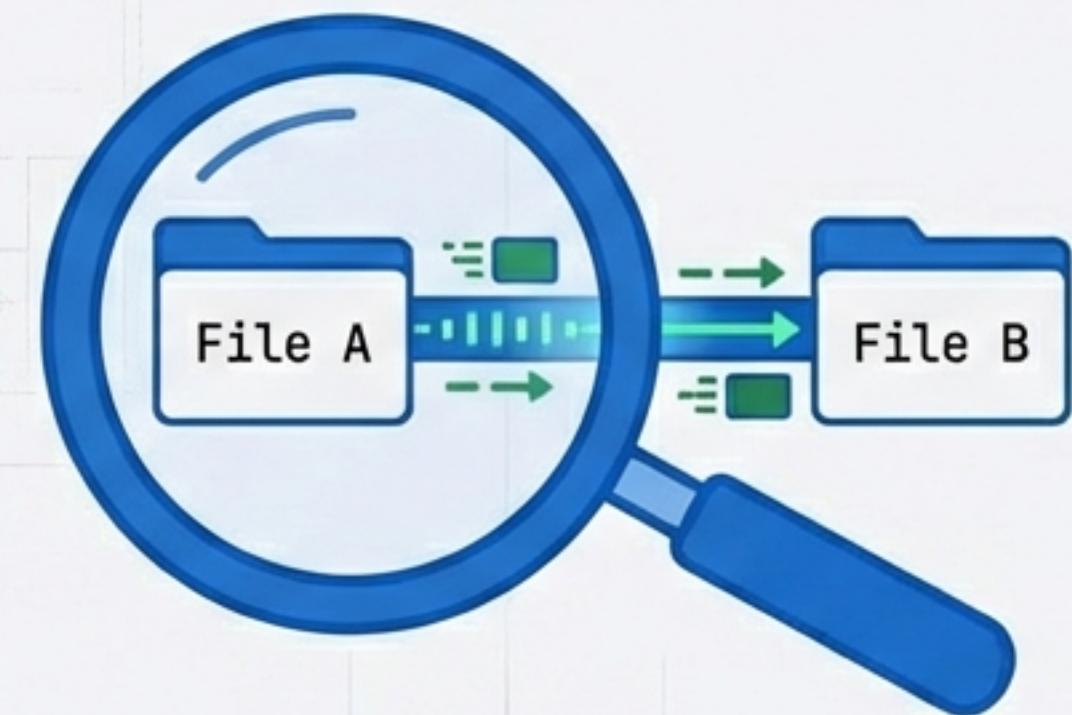


Summary: The Evolution of Necessity

Databases first solved safe shared access to data.

Indexes were added to make searching fast.

Relational models were formalized to remove duplication.



Without the database process, it's just unconnected files. The Database IS the connection.